

# Transfer Learning for Reinforcement Learning.



---

Yohannis K Telila  
y.telila@studenti.unipi.it  
Sept. 16

# Why Transfer Learning for DRL?

- Performance issue.
- Large amount of time and interaction with env't it take to learn suitable policy.
- Most Deep RL algorithm performance are measured primarily in a simulation.
- Solving the problem from scratch is not a natural approach
- Cost of failure

# Proposed papers

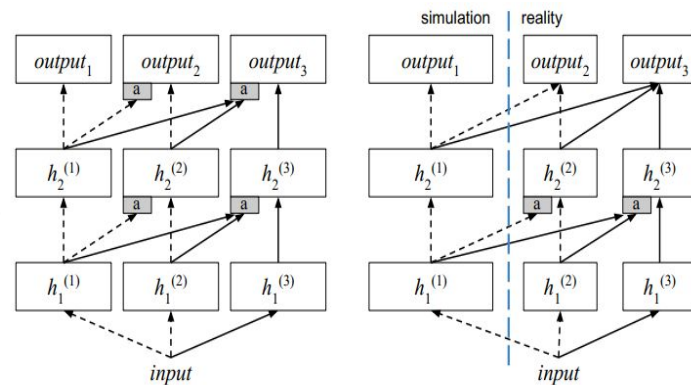
- Sim-to-Real Robot Learning from Pixels with Progressive Nets.  
Andrei et al., 2018
- Domain Adaptation For Reinforcement Learning On The Atari.  
Thomas et al., 2018
- Transfer Learning for Reinforcement Learning on a Physical Robot.  
Samuel et al., 2010.

# Sim-to-Real Robot Learning from Pixels with Progressive Nets.

Andrei et al., 2018

# Introduction

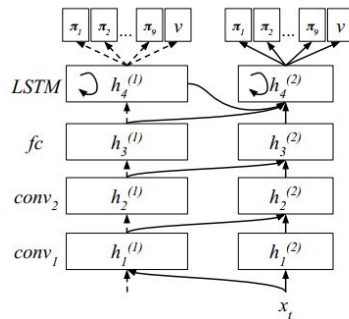




- This approach relies on the **progressive nets** architecture
- A progressive network starts with a single column.
- Columns in progressive networks are free to reuse, modify or ignore previously learned features via the **lateral connections**.
- Each column is trained to solve a particular Markov Decision Process (MDP)
- Why?
  - Feature transfer without destruction from fine tune.
  - Columns may be heterogeneous. Important for solving different task.



$$h_i^{(k)} = f \left( W_i^{(k)} h_{i-1}^{(k)} + \sum_{j < k} U_i^{(k:j)} h_{i-1}^{(j)} \right),$$








# Experiment

- The task of reaching to a visual target with Jaco arm (MuJoCo physics simulator).
- Reward of +1 if its palm is within 10cm of the target. And 50 step of episodes.
- Input : 3x64x64 pixels; output: 28 (9 discrete joint policies plus one value function).



	feedforward		recurrent	
	wide	narrow	wide	narrow
fc (output)	28	28	28	28
LSTM	-	-	128	16
fc	512	32	128	16
conv 2	32	8	32	8
conv 1	16	8	16	8
params	621K	39K	299K	37K

# Result



# Domain Adaptation For Reinforcement Learning

Thomas et al., 2018



# Introduction.

- This approach relies on the **Domain Adaptation technique**.
- One way to view learning within DRL is to consider the hidden layers as learning a state representation on top of which a policy can be learned.
- **Domain Adaptation** methods in this context are about learning to construct that state space for corresponding inputs in the target domain
- The proposed approach uses **Adversarial Auto-encoder(AAE)** to impose this regularisation in an unsupervised manner. (Adversarial Domain Adaptation)
- If the tasks are similar we expect similar embeddings. [This is in pre training step]




Figure 2: Agent's architecture for domain adaptation

# Experiment.

- Experiment carried out on three games; Pong, Breakout and Tennis to verify the effectiveness of the approach.[somehow related]
- Experiment procedure.
  1. Train agent to play Pong in the standard RL fashion;[A2C]
  2. Run the agent through a few games, capturing the final hidden layer output. [100k samples collected]
  3. Using these two data sets train the AAE.
  4. Take the weights of the Generator of our AAE, to initialise the weights of a new Policy network for the target task.
  5. Train the policy on the target task.




Fig. 1: Pong (left) to Breakout (middle) and Tennis (right)

# Result.

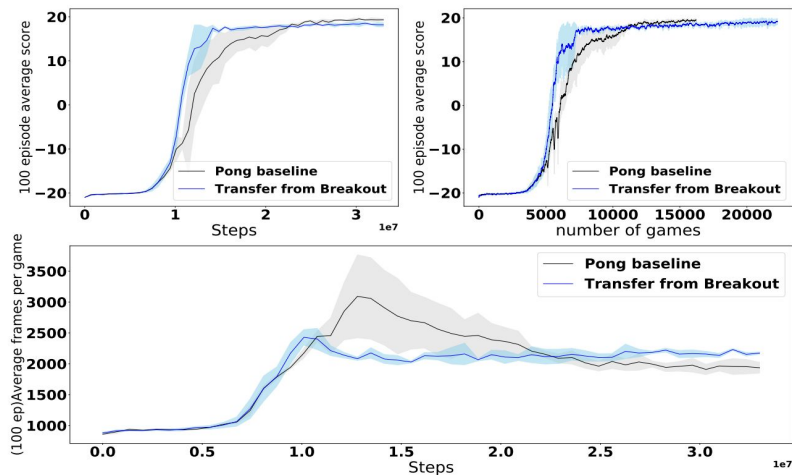


Fig: **TopLeft:** Learning Curve for Pong, **TopRight:** Plots the Average score against number of games completed. **Bottom:** 100 episode average number of frames in a game.

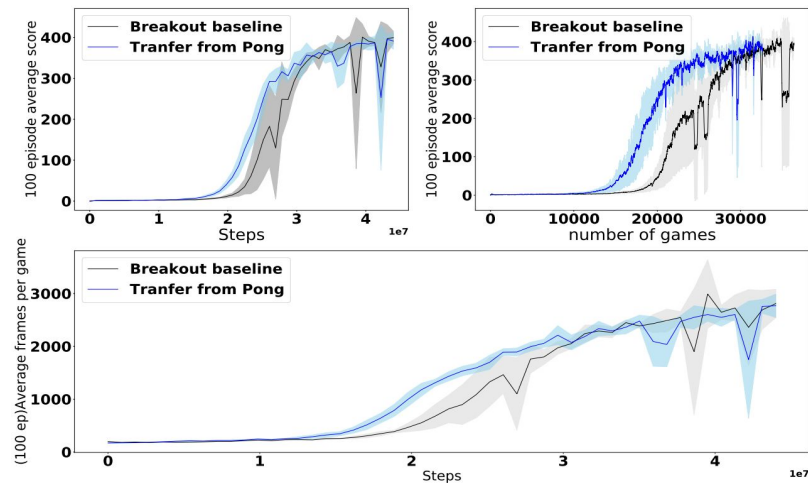


Fig: **TopLeft:** Learning Curve for Breakout, **TopRight:** Plots the Average score against number of games completed. **Bottom:** 100 episode average number of frames in a game

# Transfer Learning for Reinforcement Learning on a Physical Robot.

Samuel et al., 2010.

# Introduction.

- This approach uses “**Q-value reuse**” with **Sarsa( $\lambda$ )** to transfer information from source task to a new target task.
- Involves reusing value function  $Q_{\text{source}}$  learned from Source task, as a starting point for a new problem value Function  $Q_{\text{target}}$
- Problem : The state and action space might not coincide.
- Soln: The agent has been given the mapping

$$\chi_X(s_{\text{target}}) = s_{\text{source}}$$

$$\chi_A(a_{\text{target}}) = a_{\text{source}}$$

- The agent's new value function is given by

$$Q(s, a) = Q_{\text{source}}(\chi_X(s), \chi_A(a)) + Q_{\text{target}}(s, a)$$

- If there is no corresponding action in the source task, they are Initialized by average action values across all states.(it's a choice)

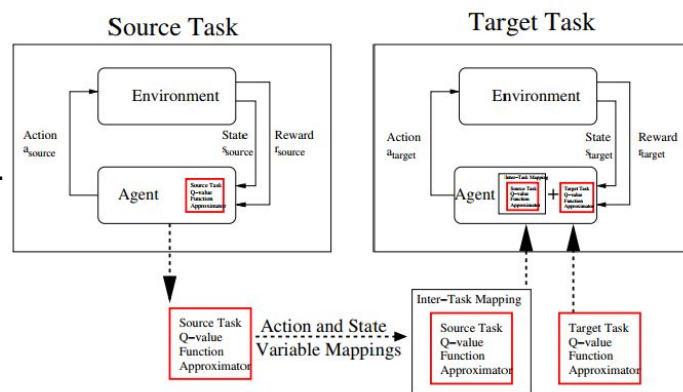



Figure 1: Q-Value Reuse

# Experiment.

- The robot's task was to hit an orange ball as far as possible at a  $45^\circ$  angle with its right hand.
- The reward signal is given by  $r = d * \cos(\theta)$ ,  
 $d$ -distance
- Source Task:
  - The robot will only use two shoulder joint and 4 observations [position and velocity of each joint].
- Target Task:
  - The robot can use all 4 joints to help him kick the ball and eight observations [the position and velocity for each joint].
- The robot will learn faster in the source task because of simpler problem.
- Both the source and target tasks were learned on the physical robot [simulation also tested].

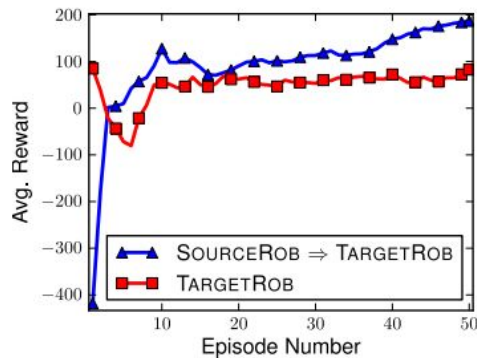


(a) Source task




(b) Target task


# Result.



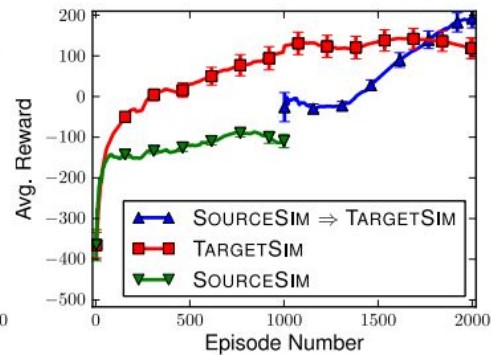
(a) Weak Transfer



(b) Strong Transfer



(a) Weak Transfer



(b) Strong Transfer

# Summary.

Transfer learning vs Domain adaptation

- Multiple tasks may be learned sequentially, without needing to specify source and target tasks.

MORE RELATED PAPERS:

- Latent Structure Matching for knowledge Transfer in Reinforcement Learning. Yi et al., 2020
- Efficient Deep Reinforcement Learning via Adaptive Policy Transfer. Tianpei et al., 2020



END.

THANK YOU !